

# B551 Assignment 6: Machine learning

Fall 2015

Due: Sunday December 13, 11:59PM

(You may submit up to 48 hours late for a 10% penalty.)

This assignment will give you a chance to implement machine learning techniques on a realistic classification problem.

We’ve assigned you a team for this assignment, based on your teamwork preferences you expressed after Assignment 5. You should only submit **one** copy of the assignment for your team, through GitHub, as described below. Both people on the team will receive the same grade on the assignment, except in very unusual circumstances. Please read the instructions below carefully; we cannot accept any submissions that do follow the instructions given here. Most importantly: please **start early**, and ask questions on the OnCourse Forum or in office hours.

**Academic integrity.** You and your partner may discuss the assignment with other people at a high level, e.g. discussing general strategies to solve the problem, talking about Python syntax and features, etc. You may also consult printed and/or online references, including books, tutorials, etc., but you must cite these materials (e.g. in source code comments). However, the work and code that you and your partner must be your group’s own work, which the two of you personally designed and wrote. You may not share written answers or code with any other students except your own partner, nor may you possess code written by another student who is not your partner, either in whole or in part, regardless of format.

In this assignment we’ll study a straightforward image classification task. These days, all modern digital cameras include a sensor that detects which way the camera is being held when a photo is taken. This metadata is then included in the image file, so that image organization programs know the correct orientation – i.e., which way is “up” in the image. But for photos scanned in from film or from older digital cameras, rotating images to be in the correct orientation must typically be done by hand.<sup>1</sup>

Your task in this assignment is to create a classifier that decides the correct orientation of a given image, as shown in Figure 1.

**Data.** To help you get started, we’ve prepared a dataset of images from the Flickr photo sharing website. The images were taken and uploaded by real users from across the world, making this a challenging task on a very realistic dataset.<sup>2</sup> Since image processing is beyond the scope of this class, we don’t expect you to implement any special techniques related to images in particular. Instead, we’ll simply treat the raw images as numerical feature vectors, on which we can then apply standard machine learning techniques. In particular, we’ll take an  $n \times m \times 3$  color image (the third dimension is because color images are stored as three separate planes – the red, green, and blue), and append all of the rows together to produce a single vector of size  $1 \times 3mn$ .

We’ve done this work for you already, so that you can treat the images simply as vectors and not have to worry about them being images at all. Your GitHub repo includes two ASCII text files, one for the training dataset and one for testing, that contain the feature vectors. To generate this file, we rescaled each image to a very tiny “micro-thumbnail” of  $8 \times 8$  pixels, resulting in an  $8 \times 8 \times 3 = 192$  dimensional feature vector. The text files have one row per image, where each row is formatted like:

```
photo_id correct_orientation r11 g11 b11 r12 g12 b12 ...
```

<sup>1</sup>and if you happen to have a new Nexus 5X smartphone, you may have to orient them by hand again too, thanks to a bug/feature: <http://www.slashgear.com/dont-panic-if-nexus-5x-camera-apps-show-inverted-images-09413669/>.

<sup>2</sup>All images in this dataset are licensed under Creative Commons, and there are no copyright issues with using them for classroom purposes under fair use guidelines. However, copyrights are still held by the photographers in most cases; you should not post these images online or use these images for any commercial purposes without checking with the original photographer, who can be identified using the Flickr URL described below.




Figure 1: Some sample images, and their correct orientation labels.

where:

- `photo_id` is a photo ID for the image.
- `correct_orientation` is 0, 90, 180, or 270. Note that some small percentage of these labels may be wrong because of noise in the data; this is just a fact of life when dealing with data from real-world sources.
- `r11` refers to the red pixel value at row 1 column 1, `r12` refers to red pixel at row 1 column 2, etc.

Although you can get away with just using the above text files, you may want to inspect the original images themselves, for debugging or analysis purposes, or if you want to change something about the way we've created the feature vectors (e.g. experiment with larger or smaller “micro-thumbnails”). You can view the images in two different ways:

- You can view the original high-resolution image on Flickr.com by taking just the numeric portion of the `photo_id` in the file above (e.g. if the `photo_id` in the file is `test/123456.jpg`, just use 12345), and then visiting the following URL:  
`http://www.flickr.com/photo_zoom.gne?id=numeric_photo_id`
- We've provided a zip file of all the images in JPEG format, available via OnCourse. We've reduced the size of each image to a  $75 \times 75$  square, which still (usually) preserves enough information to figure out the image orientation. The ZIP file also includes UNIX scripts that will convert images in the zip file to the ASCII feature vector file format above. If you want, this lets you experiment with modifying the script to produce other feature vectors (e.g. smaller sized images, or in different color mappings, etc.) and to run your classifier on your own images.

The training dataset consists of about 10,000 images, while the test set contains about 1,000. For the training set, we've rotated each image 4 times to give four times as much training data, so there are about 40,000 lines in the `train.txt` file (the training images on Flickr and the ZIP file are all oriented correctly already). In `test.txt`, each image occurs just once (rotated to a random orientation) so there are only about 1,000 lines. If you view the test images on Flickr, they'll be oriented correctly, but the ones in the ZIP file are randomly oriented.

**Step 0: Getting started.** You can find your assigned teammate by logging into IU Github, at `http://github.iu.edu/`. In the upper left hand corner of the screen, you should see a pull-down menu. Select `cs-b551`. Then in the yellow box to the right, you should see a repository called `userid1-userid2-a6`, where the other user ID corresponds to your teammate. To get started, clone the github repository:

```
git clone https://github.iu.edu/cs-b551/your-repo-name-a6
```

where *your-repo-name* is the one you found on the GitHub website above.

There's no skeleton code this time – you're free to write your code anyway you'd like, as long as it conforms to the specifications below.

**Step 1: *k*-nearest neighbor.** Implement a *k*-nearest neighbor classifier to predict the label of each test image. Write a program that is run from the command line like this:

```
python orient.py train_file.txt test_file.txt knn k
```

and then trains the *k*-nearest neighbor classifier using the specified training file, runs the classifier on the specified test file using the given value of *k*, and displays the classification accuracy (in terms of percentage of correctly-classified images). It should also display a confusion matrix and should output a file called `knn_output.txt` which indicates the estimated label for each image in the test file. The confusion matrix is a table with four rows and four columns. The entry at cell *i, j* in the table should show the number of test exemplars whose correct label is *i*, but that were classified as *j* (e.g. a “10” in the second column of the third row would mean that 10 test images were actually oriented at 180 degrees, but were incorrectly classified by your classifier as being at 180 degrees). A perfect classifier will have a diagonal confusion matrix, but we're unlikely to achieve that here. The `knn_output.txt` file should correspond to one test image per line, with the `photo_id`, a space, and then the estimated label, e.g.:

```
test/124567.jpg 180
test/8234732.jpg 0
```

**Step 2: *Neural network classification.*** Implement a fully-connected feed-forward network to classify image orientation, and implement the backpropagation algorithm to train the network using gradient descent. Your network should have one hidden layer (i.e. three layers total – the input layer, the hidden layer, and the output layer). Your program should be run like this:

```
python orient.py train_file.txt test_file.txt nnet hidden_count
```

and should train on the training file, using the specified number of hidden nodes, test on the testing file, and display the classification accuracy, confusion matrix, and output `nnet_output.txt`, with the same format as in step 1.

**Step 3: *Analysis and improvement.*** Each of the above machine learning techniques has a number of parameters and design decisions. For example, *k*-nearest neighbors has *k* and a distance function that you need to choose. Neural networks have network structure parameters (number of layers, number of nodes per hidden layer), as well as which activation function to use, whether to use traditional or stochastic gradient descent, which learning rate to use, etc. It would be impossible to try all possible combinations of all of these parameters, so identify a few parameters and conduct experiments to study their effect on the final classification accuracy. In your report, present neatly-organized tables or graphs showing classification accuracies and running times as a function of the parameters you choose. Which classifiers and which parameters would you recommend to a potential client? How does performance vary depending on the training dataset size, i.e. if you use just a fraction of the training data? Show a few sample images that were classified correctly and incorrectly. Do you see any patterns to the errors?

Finally, modify your code so that when run like this:

```
python orient.py train_file.txt test_file.txt best model_file
```

it uses whichever algorithm and parameter settings you recommend to give the best accuracy. The `model_file` is an optional parameter that you can choose to implement; if you'd like, your training routines can write their learned model to disk, so that when we run the “best” algorithm, training does not actually occur but instead a pre-trained `model_file` is loaded. Please use this option if your models take a long time to train

(more than a few minutes). As in the past, a small percentage of your assignment grade will be based on how accurate your “best” algorithm is with respect to the rest of the class. We will use a separate test dataset, so make sure to avoid overfitting!

### What to turn in

Turn in your source code file(s) as well as a **PDF document** containing your report for Step 3 above by pushing to GitHub (remember to **add**, **commit**, **push**) — we’ll grade whatever version you’ve put there as of 11:59PM on the due date. To make sure that the latest version of your work has been accepted by GitHub, you can log into the [github.iu.edu](https://github.iu.edu) website and browse the code online.

In addition to the information required in Step 3, either your report or the comments at the top of your source code should include a brief description of how your code works, and a discussion of any problems you faced, any assumptions, simplifications, and/or design decisions you made.